



TECHNICAL REPORT · LAB 01

House Price Prediction using Linear Regression

A reproducible comparative regression study across the Ames, California, and UCI Real Estate datasets — from a transparent linear baseline to regularized, tree-based, and ensemble models.

Course: MDI3003 — Advanced Predictive Analytics

Registration No.: 23MID0043

Instructor: Dr. Durgesh Kumar, SCOPE, VIT Vellore

Date: July 2026 · **Random seed:** 42 · **Python:** 3.11



Contents

1. Executive Summary	3
2. Business Problem	3
3. Dataset Description	4
4. Methodology	4
5. Exploratory Data Analysis and Feature Engineering	5
6. Model Development	7
7. Evaluation Results	7
8. Interpretation	11
9. Limitations and Risks	11
10. Future Improvements	12
11. Conclusion	12
Appendix A. Environment, Artifacts and Reproducibility	13
References	13

1. Executive Summary

This study developed and evaluated regression models for predicting residential property value on three publicly available datasets of increasing complexity: **Ames Housing** (1,460 individual home sales in Ames, Iowa; target `SalePrice` in USD), **California Housing** (20,640 census block groups from the 1990 U.S. Census; target median district value in USD 100,000s), and **UCI Real Estate Valuation** (414 transactions in New Taipei City, Taiwan; target price per unit area in 10,000 NTD per ping). Seven models — Linear Regression, Ridge, Lasso, Elastic Net, Decision Tree, Random Forest, and Gradient Boosting — were compared inside identical leakage-safe scikit-learn pipelines, with 5-fold cross-validation on the training split used for model selection and a single held-out 20% test split used once for final evidence.

Ensemble methods achieved the best generalization on all three datasets. On Ames, the selected **Gradient Boosting** model reached a test MAE of **\$16,556**, RMSE of **\$26,504**, and R^2 of **0.908** — a 69.8% RMSE reduction relative to a mean-prediction baseline. On California, **Random Forest** reached MAE 0.327, RMSE 0.505 (USD 100k), and R^2 0.805 (–55.9% vs. baseline). On UCI, **Random Forest** reached MAE 3.913, RMSE 5.669, and R^2 0.808 (–56.2% vs. baseline). The strongest predictive signals were living area, overall quality, and neighborhood (Ames); median income and geography (California); and MRT-station distance and convenience-store count (UCI). The multiple linear regression baseline remained surprisingly competitive on Ames (test R^2 0.887), confirming its value as an interpretability anchor, while its collapse on California (R^2 0.576 vs. 0.805) quantifies the cost of ignoring nonlinear structure.

The models are suitable as decision-support estimates within their original data coverage. They should **not** be used for mortgage underwriting, taxation, or other high-stakes decisions without current local data, uncertainty estimates, segment-level validation, and human review.

2. Business Problem

A real-estate analytics team needs a repeatable method for estimating the market value of residential property from measurable attributes — size, age, quality, location, and nearby amenities. The learning task is supervised regression: estimate a function $\hat{y} = f(x)$ mapping property attributes to a continuous price target, evaluated on observations not used during fitting.

2.1 Stakeholders and error costs

- **Buyers and sellers** use estimates for price discovery; large errors distort affordability decisions.
- **Banks and mortgage firms** use valuation checks for collateral risk; systematic over-valuation increases credit risk, so RMSE (which penalizes large errors) is the primary metric alongside MAE.
- **Developers and investors** use estimates for screening; bias across locations can misallocate capital.
- **Public agencies** require transparent, consistent mass appraisal — favoring interpretable baselines as a reference.

2.2 Prediction units and success criteria

The three targets are deliberately different quantities and are never compared directly: Ames predicts an **individual sale price in USD**; California predicts a **median block-group value** in USD 100,000s (not an individual house price); UCI predicts **price per unit area** in 10,000 NTD per ping (not total price). Success criteria: (i) materially lower CV/test RMSE and MAE than a naive mean-prediction baseline; (ii) an acceptable gap between cross-validation and test performance; (iii) no unexplained severe residual pattern or leakage; (iv) errors expressed in meaningful currency units.

3. Dataset Description

Table 1 summarizes the three datasets. They form a progression from a small clean problem (UCI) through a large numeric problem (California) to a realistic high-dimensional mixed-type problem (Ames), which is the primary extended project.

Table 1. Dataset summary

Dataset	Rows	Predictors	Types	Target and unit	Source / coverage
Ames Housing	1,460	80 (37 num, 43 cat, excl. Id)	Mixed	SalePrice, USD	Ames, Iowa residential sales, 2006–2010 (De Cock, 2011) ^[1]
California Housing	20,640	8 numeric	Numeric	MedHouseVal, USD 100k (block-group median)	1990 U.S. Census, via scikit-learn ^{[2][3]}
UCI Real Estate	414	6 numeric	Numeric	Price per unit area, 10k NTD/ping	New Taipei City, Taiwan, 2012–2013 (UCI ID 477) ^[4]

Data-quality audit findings. Ames contains no duplicate rows, but 19 of 81 columns have missing values; the highest rates are structural absence codes rather than unknown data (e.g., `PoolQC` 99.5% missing means “no pool”, `Alley` 93.8% means “no alley access”) and were treated as a distinct category rather than blindly imputed. The target is strongly right-skewed (skew 1.88; mean \$180,921 vs. median \$163,000; range \$34,900–\$755,000). California has no missing values but a censored target (median values capped at \$500,001) and block-group aggregation. UCI is small, complete, and clean — six numeric predictors with no missing values.

4. Methodology

4.1 Experimental protocol

All three datasets were processed with the identical protocol, making cross-dataset comparison fair:

- **Holdout design.** 80/20 random train–test split with `random_state=42` before any learned preprocessing (Ames: 1,168/292; California: 16,512/4,128; UCI: 331/83).
- **Leakage-safe preprocessing.** A `ColumnTransformer` inside a scikit-learn `Pipeline`: median imputation + standard scaling for numeric features; most-frequent imputation + one-hot encoding (`handle_unknown='ignore'`) for categorical features. All transformers are fit on training folds only.

- **Model selection.** 5-fold cross-validation (shuffled, seed 42) on the training split only. The test set was used exactly once for final evaluation of all candidates.
- **Tuning.** `GridSearchCV` over Ridge $\alpha \in [10^{-3}, 10^3]$ (13 log-spaced values), scored by negative CV RMSE.
- **Reproducibility.** Fixed seed 42 everywhere; recorded package versions (Appendix A); serialized final pipelines with `joblib`; raw data unchanged.

4.2 Models and theory anchor

Ordinary Least Squares estimates coefficients by minimizing the sum of squared residuals:

$$\hat{\beta} = \arg \min_{\beta} \sum_i (y_i - \beta_0 - \beta^T x_i)^2 = (X^T X)^{-1} X^T y \quad (1)$$

Multiple Linear Regression is the scientific baseline and interpretability anchor of the study. Ridge adds an L2 penalty $\lambda \sum \beta_j^2$ to stabilize correlated coefficients; Lasso adds an L1 penalty $\lambda \sum |\beta_j|$ that can zero out coefficients; Elastic Net mixes both. Decision Trees capture nonlinearity and interactions but are high-variance; Random Forest reduces that variance by bagging; Gradient Boosting corrects errors sequentially. Predictive accuracy is reported as MAE and RMSE in original target units, plus the coefficient of determination:

$$R^2 = 1 - \Sigma(y_i - \hat{y}_i)^2 / \Sigma(y_i - \bar{y})^2 \quad (2)$$

4.3 Naive baselines

A mean-prediction baseline establishes whether any learning adds value. Test-set baseline RMSE is \$87,619 (Ames), 1.145 in USD 100k units (California), and 12.95 in 10k-NTD/ping units (UCI, derived exactly from the reported Linear Regression R^2). Every fitted model beats its baseline by a wide margin on Ames and California; on UCI all models except the Decision Tree improve the baseline by roughly 37–56%.

5. Exploratory Data Analysis and Feature Engineering

EDA was performed on training data only; the test set was never used for feature-engineering decisions. Four observations materially influenced modeling:

- **Right-skewed target (Ames).** Figure 1 shows the long right tail of sale prices. This justifies RMSE-sensitive evaluation and motivates the log-target extension discussed in Section 10.
- **Informative missingness (Ames).** Figure 2 shows that the most-missing features encode absence (no pool, no alley, no fence, no fireplace). Dropping these rows would discard ~90% of the data; they were encoded as explicit “None” categories instead.
- **Strong but correlated size/quality signals.** Figure 3 ranks the top numerical correlates of price — OverallQual ($r = 0.79$) and GrLivArea ($r = 0.71$) dominate, while garage size/count and basement/first-floor area are mutually redundant, motivating regularized models.
- **Nonlinear geography.** California price surfaces are strongly driven by latitude/longitude and median income with clear saturation, and UCI price falls non-linearly with MRT distance — patterns a purely linear model cannot represent (confirmed in Section 7).




Figure 1. Ames Housing: sale-price distribution (n = 1,460). Right skew 1.88; mean exceeds median by ~\$18k.




Figure 2. Ames Housing: top-10 features by missing-value rate. High rates reflect structural absence, not data corruption.




Figure 3. Ames Housing: strongest numerical correlates of SalePrice (Pearson r, full dataset).

6. Model Development

The experiment matrix followed the laboratory design: a dummy mean baseline (E0), a simple one-feature Linear Regression (E1: GrLivArea for Ames, MedInc for California, MRT distance for UCI), multiple Linear Regression through the full pipeline (E2), regularized variants (E3), a depth-limited Decision Tree and 300-tree Random Forest (E4), Gradient Boosting (E5), repetition on all three datasets (E6), and Ridge α tuning via grid search as the controlled tuning experiment (E7).

Key hyperparameters: Decision Tree `max_depth=6`; Random Forest `n_estimators=300`, `min_samples_leaf=2`; Gradient Boosting with scikit-learn defaults; Lasso/Elastic Net at $\alpha = 0.001$ with `max_iter=20000`. All models shared the same preprocessing pipeline and split, so performance differences are attributable to the estimators rather than to data handling.

7. Evaluation Results

7.1 Ames Housing (target: SalePrice, USD)

Gradient Boosting was selected on cross-validated training performance (CV RMSE \$27,709, best of seven) and confirmed on the held-out test set (Table 2, Figure 4). The linear family clusters tightly at $R^2 \approx 0.88$ – 0.90 , while the single Decision Tree clearly underfits relative to ensembles.

Table 2. Ames Housing — test-set comparison (292 held-out sales) and 5-fold CV on training data

Model	Test MAE (\$)	Test RMSE (\$)	Test R^2	CV RMSE (\$)	CV R^2
Gradient Boosting	16,555.67	26,503.77	0.9084	27,708.84	0.8668
Lasso	18,010.80	28,337.93	0.8953	36,042.14	0.7747
Random Forest	17,473.41	28,747.70	0.8923	30,637.61	0.8400
Linear Regression	18,287.70	29,473.87	0.8867	36,396.35	0.7700
Elastic Net	18,809.63	29,709.27	0.8849	33,588.56	0.8057
Ridge	19,004.41	29,841.82	0.8839	32,817.04	0.8147
Decision Tree	26,729.14	39,928.28	0.7922	44,010.85	0.6716




Figure 4. Ames Housing: test vs 5-fold CV RMSE. Gradient Boosting leads on both; the CV–test gap is small, indicating stable selection.

7.2 California Housing (target: median block-group value, USD 100k)

Random Forest was selected by CV and confirmed on test (Table 3, Figure 5). The striking finding is the linear family's collapse: Linear Regression and Ridge plateau at $R^2 \approx 0.576$, and Lasso at $\alpha = 0.001$ shrinks the model to the mean ($R^2 \approx 0.000$), while ensembles exceed $R^2 0.80$ — direct evidence that the income–price and geography–price relationships are strongly nonlinear.

Table 3. California Housing — test-set comparison (4,128 held-out block groups) and 5-fold CV (units: USD 100,000)

Model	Test MAE	Test RMSE	Test R^2	CV RMSE	CV R^2
Random Forest	0.3274	0.5051	0.8053	0.5109	0.8047
Gradient Boosting	0.3717	0.5422	0.7756	0.5321	0.7881
Decision Tree	0.4539	0.7028	0.6230	0.7333	0.5973
Ridge	0.5332	0.7456	0.5758	0.7205	0.6115
Linear Regression	0.5332	0.7456	0.5758	0.7205	0.6115
Elastic Net	0.8060	1.0219	0.2031	1.0288	0.2080
Lasso	0.9061	1.1449	−0.0002	1.1562	−0.0002




Figure 5. California Housing: test vs CV RMSE. Ensembles dominate; Lasso/Elastic Net collapse to the mean baseline at $\alpha = 0.001$.

7.3 UCI Real Estate Valuation (target: price per unit area, 10k NTD/ping)

Random Forest was again selected by CV (Table 4, Figure 6). Gradient Boosting achieved a marginally lower test MAE (3.904 vs. 3.913) but a higher test RMSE and clearly worse CV scores — under the stated selection rule (choose by cross-validated training performance, evaluate once on test), Random Forest remains the justified choice.

Table 4. UCI Real Estate — test-set comparison (83 held-out transactions) and 5-fold CV (units: 10,000 NTD per ping)

Model	Test MAE	Test RMSE	Test R ²	CV RMSE	CV R ²
Random Forest	3.9131	5.6689	0.8084	7.8302	0.6643
Gradient Boosting	3.9036	5.8443	0.7964	8.3502	0.6195
Ridge	5.3022	7.3108	0.6814	9.1083	0.5481
Elastic Net	5.3048	7.3139	0.6811	9.1105	0.5478
Lasso	5.3052	7.3144	0.6811	9.1108	0.5478
Linear Regression	5.3054	7.3148	0.6811	9.1109	0.5478
Decision Tree	5.9313	8.1530	0.6038	10.7388	0.3561




Figure 6. UCI Real Estate: test vs CV RMSE. The linear family is nearly identical across all four variants on this small dataset.

7.4 Cross-dataset synthesis

Figure 7 summarizes the selected models. Three consistent patterns emerge: (i) ensembles improve most where relationships are nonlinear and interactions matter (California: +0.23 R^2 over linear; UCI: +0.13; Ames: +0.02); (ii) the CV–test gap is small on the two larger datasets, indicating stable selection; (iii) on the 414-row UCI dataset the test R^2 (0.808) is notably higher than CV R^2 (0.664), a small-sample artifact — the CV estimate is the more conservative and trustworthy expectation of future performance.




Figure 7. Selected models: test vs cross-validation R^2 across the three datasets.

Adjusted R^2 caveat (Ames). The reported adjusted R^2 values for Ames are large negative numbers. This is an artifact, not a model failure: after one-hot encoding, the design matrix has more features (p) than test observations ($n = 292$), so the penalty term $(n-1)/(n-p-1)$ becomes negative. Adjusted R^2 is only meaningful when $n \gg p$; cross-validation is the appropriate comparison tool in the high-dimensional regime.

8. Interpretation

8.1 What drives price

Findings are stated as conditional associations, not causal effects:

- **Ames.** Above-ground living area (GrLivArea), overall material/finish quality (OverallQual), neighborhood, and garage capacity dominate both linear coefficients and tree importances. The linear baseline's test R^2 of 0.887 shows these attributes explain most of the price variance additively; Gradient Boosting's extra ~ 2 R^2 points come from interactions (e.g., quality $\times$ area) and diminishing returns at the high end.
- **California.** Median income is by far the strongest signal, followed by location (latitude/longitude, i.e., coastal proximity), average rooms, and house age. The 23-point R^2 gap between linear and ensemble models shows these effects are strongly nonlinear and interacting.
- **UCI.** Distance to the nearest MRT station, number of nearby convenience stores, geographic coordinates, and house age drive unit prices — a transit-accessibility premium consistent with urban economic intuition.

8.2 The linear baseline's role

Linear Regression never wins on accuracy, but it is far from useless. On Ames it is within $\sim \$3,000$ RMSE of the best model while remaining fully coefficient-interpretable; on UCI it matches the regularized variants almost exactly, confirming low multicollinearity there; and on California its failure is itself diagnostic, revealing how much structure is nonlinear. This is exactly the “interpretability anchor” role the study design intended.

8.3 Regularization behavior

Ridge never degraded materially below plain OLS and improved Ames CV RMSE by $\sim \$3,600$, showing mild variance reduction where one-hot categories create correlated columns. Lasso and Elastic Net behaved as expected on Ames and UCI but collapsed to the mean predictor on California at $\alpha = 0.001$ — a reminder that the useful regularization range is dataset-dependent and must be tuned by CV rather than assumed.

8.4 Generalization behavior

For both selected models the test metrics sit slightly better than or close to CV metrics (Ames: test RMSE \$26,504 vs. CV \$27,709; California: 0.505 vs. 0.511), so there is no evidence of overfitting at the selection stage. UCI shows the opposite direction (test better than CV), which is benign: it reflects fold-to-fold variance on 331 training rows rather than leakage, and argues for reporting CV uncertainty on small datasets.

9. Limitations and Risks

- **Historical drift.** All models are trained on past markets (Ames 2006–2010, California 1990, Taiwan 2012–2013). Inflation, interest-rate cycles, and market regime changes will degrade absolute price predictions; the models estimate relative structure better than current price levels.

- **Target semantics.** California predicts a *median block-group* value with censoring at \$500,001, not an individual home price; UCI predicts price *per unit area*, not total price. Using either as an individual-property valuation would be a category error.
- **Random split realism.** A random holdout assumes deployment on similar properties from the same market and period. Production use would require temporal or spatial (grouped) splits to measure true forward-looking performance.
- **Small-sample uncertainty (UCI).** With 83 test rows, a handful of transactions moves R^2 by several points; the CV estimate (0.664) should anchor expectations.
- **Fairness and governance.** Neighborhood and coordinate features can encode socioeconomic and historical inequalities. Use in credit, taxation, or access decisions requires governance review, segment-level error analysis, and human escalation rules for unusual properties.
- **Unmeasured drivers.** No dataset includes financing conditions, renovation history beyond year-remodeled, school quality, or micro-location attributes, capping attainable accuracy.
- **Limited tuning budget.** Only Ridge received a full grid search. The selected ensembles used fixed, sensible configurations; further tuning (learning rate, depth, subsampling) would likely close the remaining gap between CV leaders.

10. Future Improvements

- **Log-transformed target.** The Ames target skew (1.88) suggests `TransformedTargetRegressor` with `log1p` could reduce the influence of the high-end tail; final errors must be reported after inverse transformation.
- **Deeper tuning of ensembles.** Randomized/Bayesian search over Gradient Boosting and Random Forest hyperparameters, with a stated compute budget.
- **Spatial and temporal validation.** Neighborhood-grouped or year-based splits to measure realistic deployment performance; spatial features (distance-to-center) for California.
- **Uncertainty quantification.** Quantile regression or conformal prediction intervals so each estimate ships with an error range, as professional practice requires.
- **Segment-level error analysis.** Systematic evaluation by price band and neighborhood to detect where errors concentrate (e.g., very expensive or very old homes).
- **Deployment packaging.** Versioned artifacts (already saved via `joblib`), input-schema validation, drift monitoring, and periodic retraining.

11. Conclusion

The study met its stated success criteria on all three datasets. Every final model reduces RMSE by 56–70% relative to a naive mean baseline, generalizes stably (small CV–test gaps on the larger datasets), and was produced by a fully reproducible, leakage-safe pipeline with a fixed seed, recorded versions, and serialized artifacts. The central analytical question — how accurately, reliably, and transparently can property value be predicted — has a dataset-dependent answer: on the rich mixed-type Ames data, ensemble boosting reaches $R^2 = 0.908$ with the linear baseline close behind and fully interpretable; on California and UCI, nonlinear structure is essential and tree

ensembles are clearly superior. Linear Regression thus proves its dual value: a competitive, transparent baseline on additive problems, and a diagnostic instrument that exposes nonlinearity when it fails. Within their data coverage, the selected models are fit for decision support; beyond it — current markets, individual California homes, or high-stakes credit decisions — they require revalidation, uncertainty estimates, and human oversight.

Appendix A. Environment, Artifacts and Reproducibility

Table 5. Execution environment (from `run_metadata.json`; executed 2026-07-19)

Component	Version / setting	Component	Version / setting
Python	3.11.0	scikit-learn	1.8.0
numpy	1.26.4	matplotlib	3.10.8
pandas	2.3.3	seaborn	0.13.2
joblib	1.5.3	Random seed	42 (all splits, CV, stochastic models)

Table 6. Submitted artifacts per dataset

Dataset	Selected model	Model artifact (joblib)	Result CSVs
Ames Housing	Gradient Boosting	<code>ames_house_price_pipeline.joblib</code>	<code>ames_model_comparison.csv</code> , <code>ames_cross_validation.csv</code>
California Housing	Random Forest	<code>california_house_price_pipeline.joblib</code>	<code>california_model_comparison.csv</code> , <code>california_cross_validation_results.csv</code>
UCI Real Estate	Random Forest	<code>uci_real_estate_model.joblib</code>	<code>uci_model_comparison.csv</code> , <code>uci_cross_validation.csv</code>

Supporting files: three self-contained analysis notebooks (Ames, California, UCI), `23MID0043_Lab01_README.md` (execution order and environment), `execute_notebooks.py` (headless notebook re-execution harness), and `run_metadata.json` (seed, split sizes, selected models, headline metrics). Notebooks restart-and-run-all cleanly; raw data is never modified.

References

1. De Cock, D. (2011). Ames, Iowa: Alternative to the Boston Housing Data as an End of Semester Regression Project. *Journal of Statistics Education*, 19(3).
2. Pedregosa, F., Varoquaux, G., Gramfort, A., Michel, V., Thirion, B., Grisel, O., ... Duchesnay, É. (2011). Scikit-learn: Machine learning in Python. *Journal of Machine Learning Research*, 12, 2825–2830.
3. Pace, R. K., & Barry, R. (1997). Sparse spatial autoregressions. *Statistics & Probability Letters*, 33(3), 291–297.
4. Yeh, I.-C., & Hsu, T.-K. (2018). Building real estate valuation models with comparative approach through case-based reasoning. *Applied Soft Computing*, 65, 260–271. (UCI Machine Learning Repository, dataset ID 477.)
5. James, G., Witten, D., Hastie, T., Tibshirani, R., & Taylor, J. (2023). *An Introduction to Statistical Learning with Applications in Python*. Springer.
6. Kumar, D. (2026). *MDI3003 — Advanced Predictive Analytics: Laboratory Instruction Manual, Lab 01*. SCOPE, VIT Vellore.